

# Lecture 2:

# Designing for Human-AI Interaction

CS698Y 2025-26 Semester-1



# Last class...

- AI is everywhere
  - Algorithm-centric → Human-centric
    - Accuracy is not enough!
    - Various examples of accurate AI model use questioned
    - Importance of human aspects
  - HAI is different from traditional user interaction design
    - Uncertainty, non-determinism, error, trust, initiative
    - Evolves over time (e.g., with more data, online learning)
    - Explanations and transparency (no one wants arbitrary black box predictions; e.g., course recommendations)
-



# Today:

- Key HAI principles / principles
  - How to design according to those principles / guidelines
  - Hands-on exercise
-

# Why do we need guidelines?

- Avoid common pitfalls (over trust, confusion, misuse)
- Create consistency in user experience
- Promote fairness, safety, usability
- Bridge the gap between developers and users
- Note: guidelines are “rules of thumb”
  - Practical, comes from experience, but need to be routinely revisited
  - New system kinds => new guidelines
  - Don't be myopic → look at experiences beyond individual features

# Key practical guidelines (from Industry)

- Google's People AI guidelines (PAIR) → incl. ethical aspects
- Microsoft's Human AI experiences (HAX) toolkit
- Google's Explainability Rubric
- HCI Institute's toolkit (for ideation)
- A small set of common needs, mistakes, pitfalls
  - Abstracted out to guidelines, patterns, etc.

# Steps

- What to build?
- Decide on level of control?
- Ensure a good user experience (HAX guidelines)?
- How to evaluate whether your AI is successful or not?

# Step 1: What to build? Do we need AI?

- Two ways features / products are built
  - People-centric: user needs
  - Technology-centric: have the hammer & look for a nail!
- Former is required for long-term success of a product
- So, keep people at its centre!

# What is AI good at? (CMU HCII AI Design Toolkit)

Detection

Identification

Estimation

Forecast /  
Predict

Compare

Discover

Generate

Act



# Step 1: What to build (contd.)

- How pervasive is the problem?
- Who does it affect?
- How does this problem affect various individuals?
- How does this problem vary across different groups and dimensions of identity?
- Does the problem's impact change over time?

# Step 2: Decide on the role of AI

- Automation vs. augmentation
  - Should AI automate things for and act on your behalf?
    - With minimal user invention?
  - Should AI augment human capabilities?
    - Add on to humans, with specific strengths than people (e.g., crunch data).
- Examples:
  - Predict rainfall and declare leave
  - Predict rainfall, and let people decide what they want to do.

# Too much automation is problematic

- Errors
- People might over trust
- People get complacent and don't check routinely
- Don't automate!
  - Simply framing makes a difference here.
  - Cancer predictor → cancer testing prioritization decision support

# Too much automation is problematic

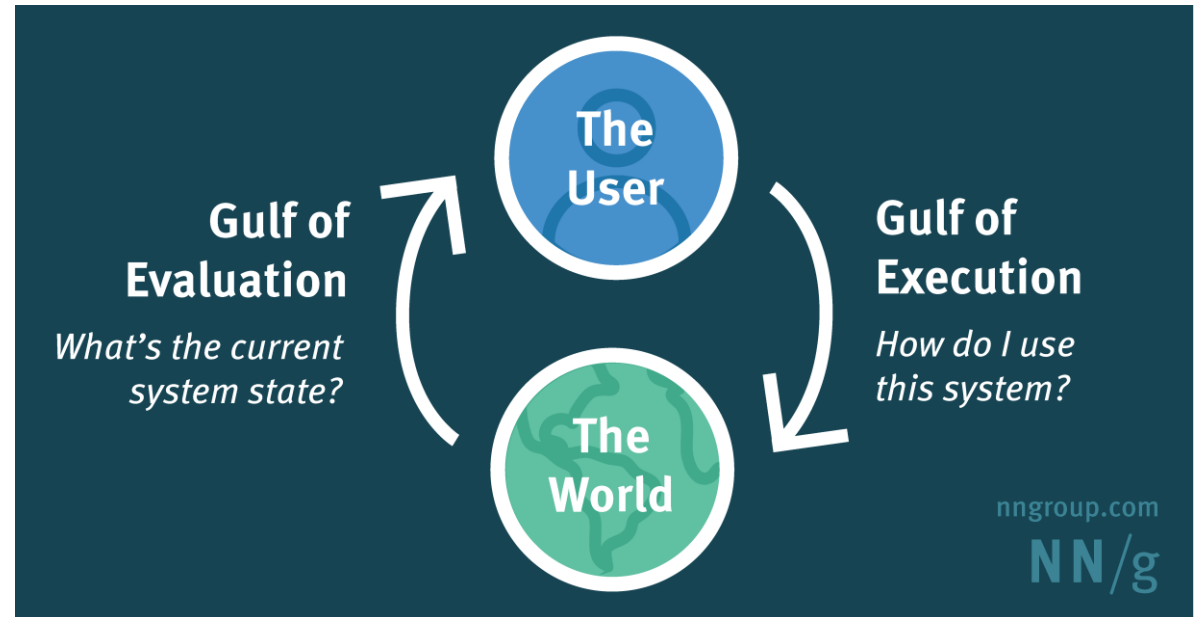
- What would they hand out to an assistant?
  - If not, then don't stick an AI.
- If yes, what skills do they expect the assistant to know?
  - If AI doesn't have such skills, don't replace human with AI / frame carefully.
  - How much oversight/supervision/trust → if none, frame carefully.
- Repetitive, too large/resource constraints, minimal oversight
  - These are great times to automate!

# Also consider initiative

- User takes initiative to invoke AI
  - Spell check turn on!
  - Ask Generative AI
- AI takes initiative, but user can ignore / accept as final action
  - (Mixed initiative): Google doc spell check
- AI takes initiative and acts on behalf of user
  - Typo replacement

- Step 1: What to build?
- Step 2: How to build?
  - Initiative, automation vs. augmentation
- Step 3: How to build well?

# Bridging Norman gulfs: Good usability



# Microsoft's HAX guidelines





# HAX guidelines Set #1: Initial context

- Why is this important?
  - Not clear how to use a system
  - Incorrect use → leading to frustration
  - Under-use of the system
- Make clear how well it can do that
  - Not lead to over trust, or under trust in system.

Ask anything

+ 🔧 Tools



ChatGPT can make mistakes. Check important info. See [Cookie Preferences](#).

## Initially

1



**Make clear  
what the  
system can  
do.**

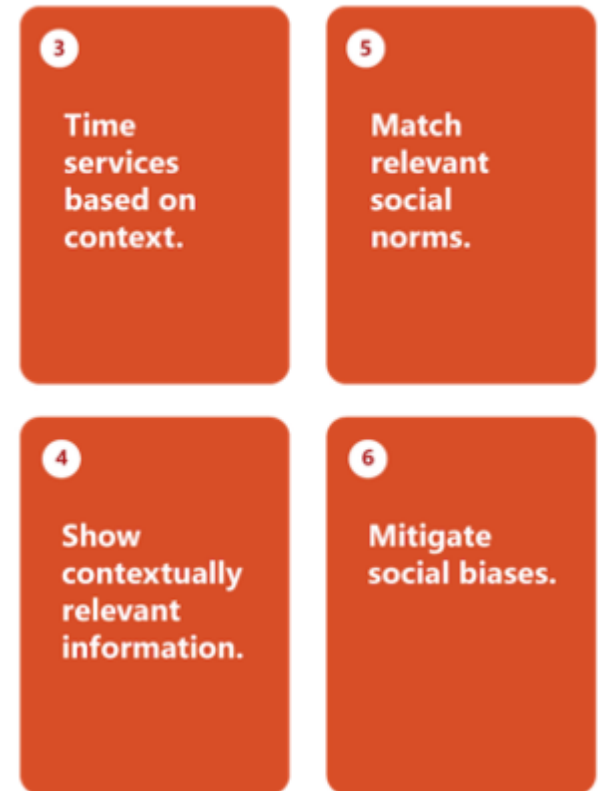
2

**Make clear  
how well  
the system  
can do what  
it can do.**

# HAX guidelines Set #2: During Interaction

- Time services based on context
  - When to provide recommendations/suggestions/...
  - Examples:
    - Google Map Navigation: “In 200 m, turn left”
    - Rewrite sentences in Word, slide design in powerpoint.
- Show contextually relevant information
  - Not a list of all possible things
  - “Also see” recommendations on Amazon.

## During interaction



# HAX guidelines Set #2: During Interaction

- Match relevant social norms
  - Look for people-people norms
  - Examples:
    - Person navigator automatically says: “in a bit, turn left”
    - India vs. Western Google Map directions
    - “Let me know if you need details” (ChatGPT)
    - Explanations—course recommendations.
- Mitigate social biases
  - Walking speed on Google Maps
  - Gender, conformity, ...

## During interaction



# HAX guidelines set #3: When wrong

- Support efficient invocation / dismissal
  - Designer in Power Point
  - If it doesn't show up, call "Designer"
  - If it shows up unwarranted, turn off.
- Support efficient correction
  - Auto spelling corrections → simply UNDO

## When wrong

- 
- 7 Support efficient invocation.
  - 9 Support efficient correction.
  - 8 Support efficient dismissal.
  - 10 Scope services when in doubt.
  - 11 Make clear why the system did what it did.

# HAX guidelines set #3: When wrong

- Scope services when in doubt
  - “Tell me about Mac”.
  - Instead of assuming, ask if it is about: Mac prefix in names or Mac computers
  - Service → scoped down, but graceful and error-proof.
- Make clear why the system did what it did
  - Google Maps → fastest, least traffic, avoids tolls.
  - Sometimes suggests silly routes, without saying why.
  - Not always easy → more on this later!

## When wrong

- 
- 7 Support efficient invocation.
  - 9 Support efficient correction.
  - 8 Support efficient dismissal.
  - 10 Scope services when in doubt.
  - 11 Make clear why the system did what it did.

# HAX guidelines set #4: Over time

- Remember recent interactions
  - Common usability need
  - Pick up where you left, recent documents...
- Learn from user behavior
  - Not done very well → privacy laws
  - “You usually open this document around this time”



# HAX guidelines set #4: Over time

- Update and adapt **cautiously**
  - Why?
  - Be transparent, notify old vs. new responses
  - Allow to go back to old version
  - Warn if an old version is going to be retired
- Encourage granular feedback



# HAX guidelines set #4: Over time

- Convey the consequences of user actions
  - Delete → say if it can be recovered or not
  - Dismiss → will not show up, until you hail from settings.
  - Internet will not be used
- Provide global controls
  - Turn on/off
  - What data to consider or not
  - Automatic vs. only when hailed ...





# HAX guidelines set #4: Over time

- Notify users about changes
  - Again, common usability guideline
  - Important, because changes are not visible
  - Clarify what changes, allow going back, ...





Initially

During interaction

When wrong

Over time

# Warning!

Do not think of individual features; but of entire experience of the user using the tool.

# How do we use these guidelines?

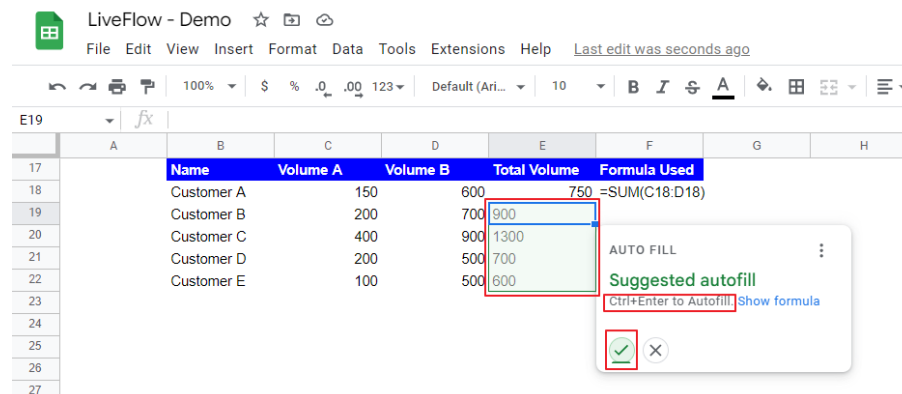
- Evaluate an existing tool / AI feature
- When designing a new AI feature
  - Easier to do it before than change it later
  - Sometimes, involves changing the model itself!

# Hands-on activity...

- ... But first, Acadly attendance!

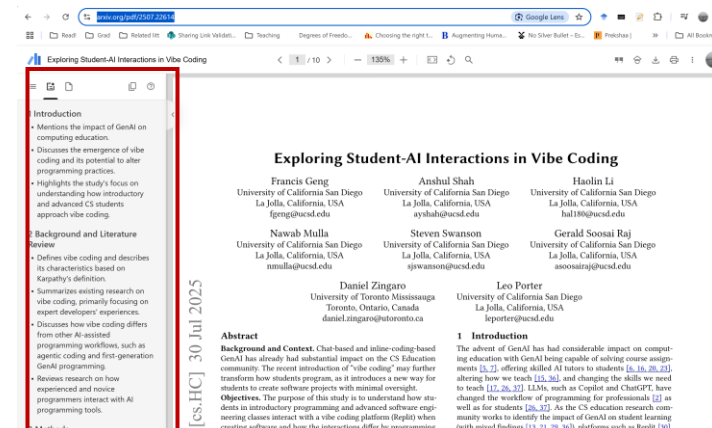
# Hands-on activity

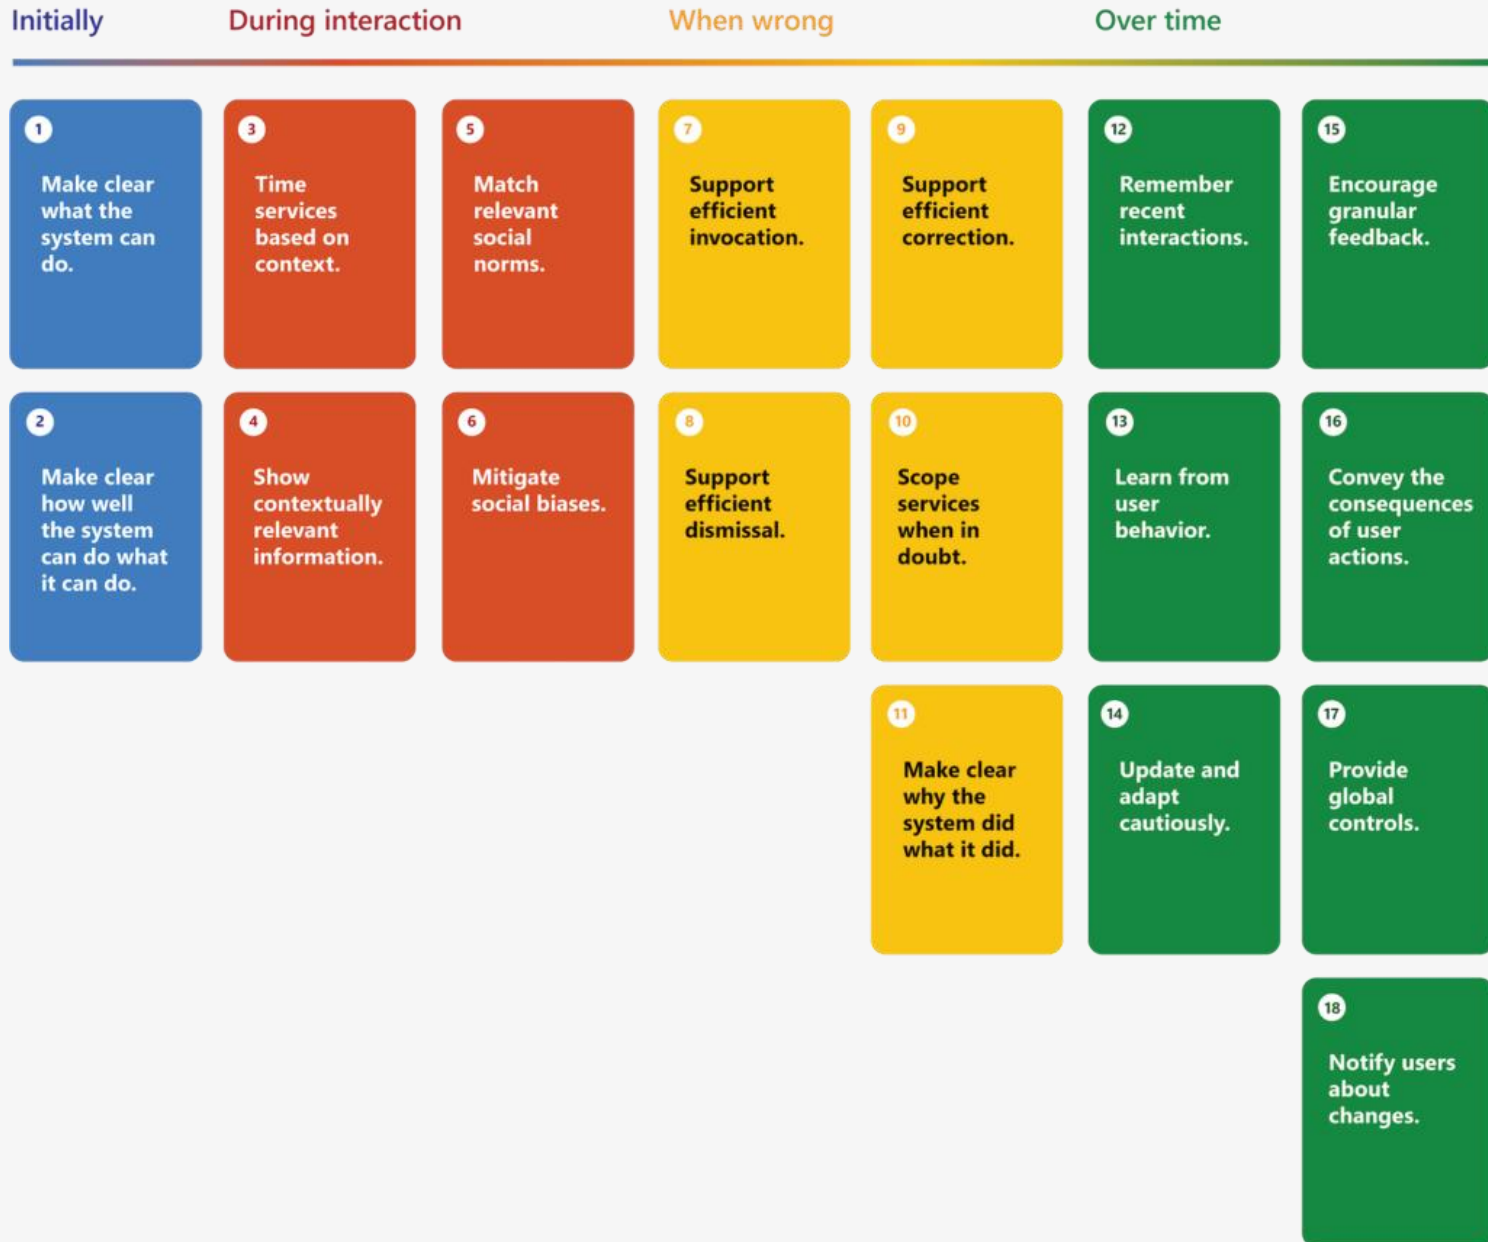
- Work in pairs
- Pick one of two features:
  - Google sheets → formula autofill
  - Google Chrome → Open PDF document → summary generated on left
- See how well the feature of your choice does on these 18 guidelines



The screenshot shows a Google Sheets interface with a table containing customer data. The table has columns for Name, Volume A, Volume B, Total Volume, and Formula Used. The 'Total Volume' column is being autofilled with the formula =SUM(C18:D18). A tooltip is visible over the cell, suggesting the autofill and providing instructions on how to accept or show the formula.

Name	Volume A	Volume B	Total Volume	Formula Used
Customer A	150	600	750	=SUM(C18:D18)
Customer B	200	700	900	
Customer C	400	900	1300	
Customer D	200	500	700	
Customer E	100	500	600	





- How do these apply to Google Sheets formula autofill / PDF document summary in Chrome?
- For each guideline:
  - Say if the guideline applies
  - If yes:
    - Does it follow the guideline? How?
    - If not, what can be done?
- Write roll number & names
- Turn in end of class

# Summary...

- Last class → importance of human-AI interaction
- Today → designing human-AI interfaces
  - Guidelines from Microsoft
  - 18, broken down to:
    - Initial
    - During interaction
    - During errors
    - Long term use
  - Tool for evaluating an interface / designing an interface

# Readings

- Google PAIR Guidebook
  - Read chapter 1: User needs + defining success
  - Read chapter 2: Mental models + expectations
- Guidelines for Human-AI Interaction, Amershi et al. 2019
  - <https://www.microsoft.com/en-us/research/wp-content/uploads/2019/01/Guidelines-for-Human-AI-Interaction-camera-ready.pdf>



Next class...